

Operating System Principles:  
Memory Management  
CS 111  
Operating Systems

# Outline

- What is memory management about?
- Memory management strategies:
  - Fixed partition strategies
  - Dynamic partitions
  - Buffer pools
  - Garbage collection
  - Memory compaction

# Memory Management

- Memory is one of the key assets used in computing
- In particular, memory abstractions that are usable from a running program
  - Which, in modern machines, typically means RAM
- We have a limited amount of it
- Lots of processes need to use it
- How do we manage it?

# Memory Management Goals

## 1. Transparency

- Process sees only its own address space
- Process is unaware memory is being shared

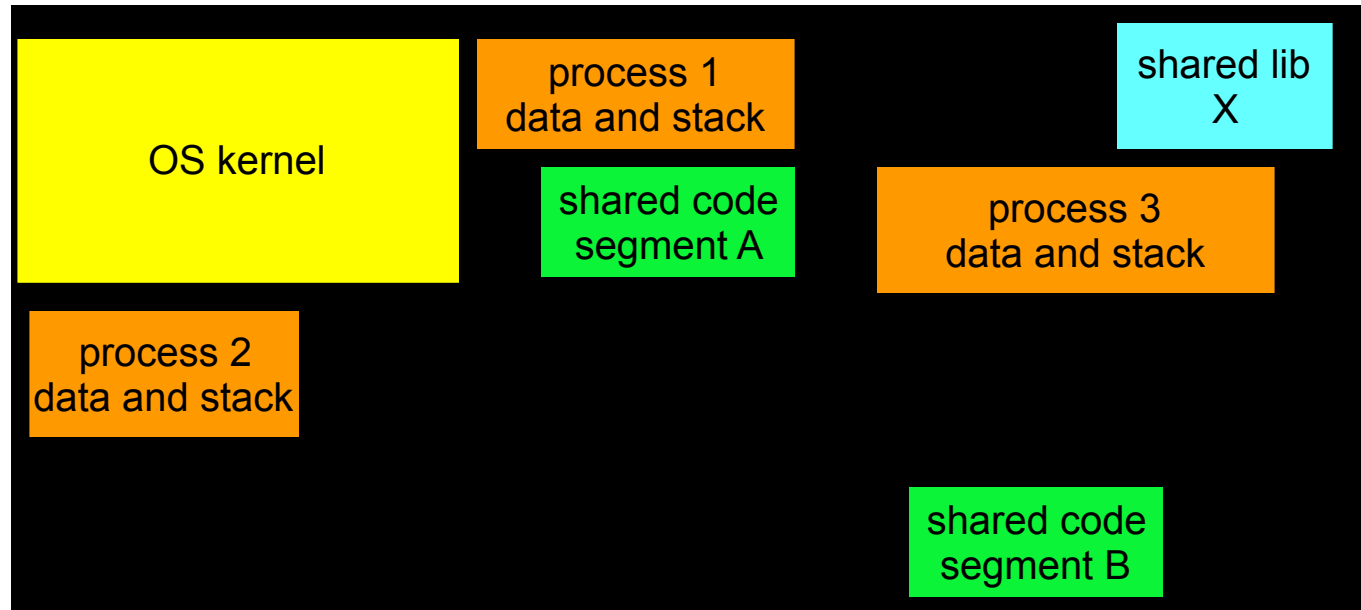
## 2. Efficiency

- High effective memory utilization
- Low run-time cost for allocation/relocation

## 3. Protection and isolation

- Private data will not be corrupted
- Private data cannot be seen by other processes

# Physical Memory Allocation



Physical memory is divided between the OS kernel, process private data, and shared code segments.

# Physical and Virtual Addresses

- A cell of RAM has a particular physical address
- Years ago, that address was used by processes to name RAM locations
- Instead, we can have processes use virtual addresses
  - Which may not be the same as physical addresses
- More flexibility in memory management, but requires virtual to physical translation

# Aspects of the Memory Management Problem

- Most processes can't perfectly predict how much memory they will use
- The processes expect to find their existing data when they need it where they left it
- The entire amount of data required by all processes may exceed amount of available physical memory
- Switching between processes must be fast
  - Can't afford much delay for copying data
- The cost of memory management itself must not be too high

# Memory Management Strategies

- Fixed partition allocations
- Dynamic partitions
- Relocation



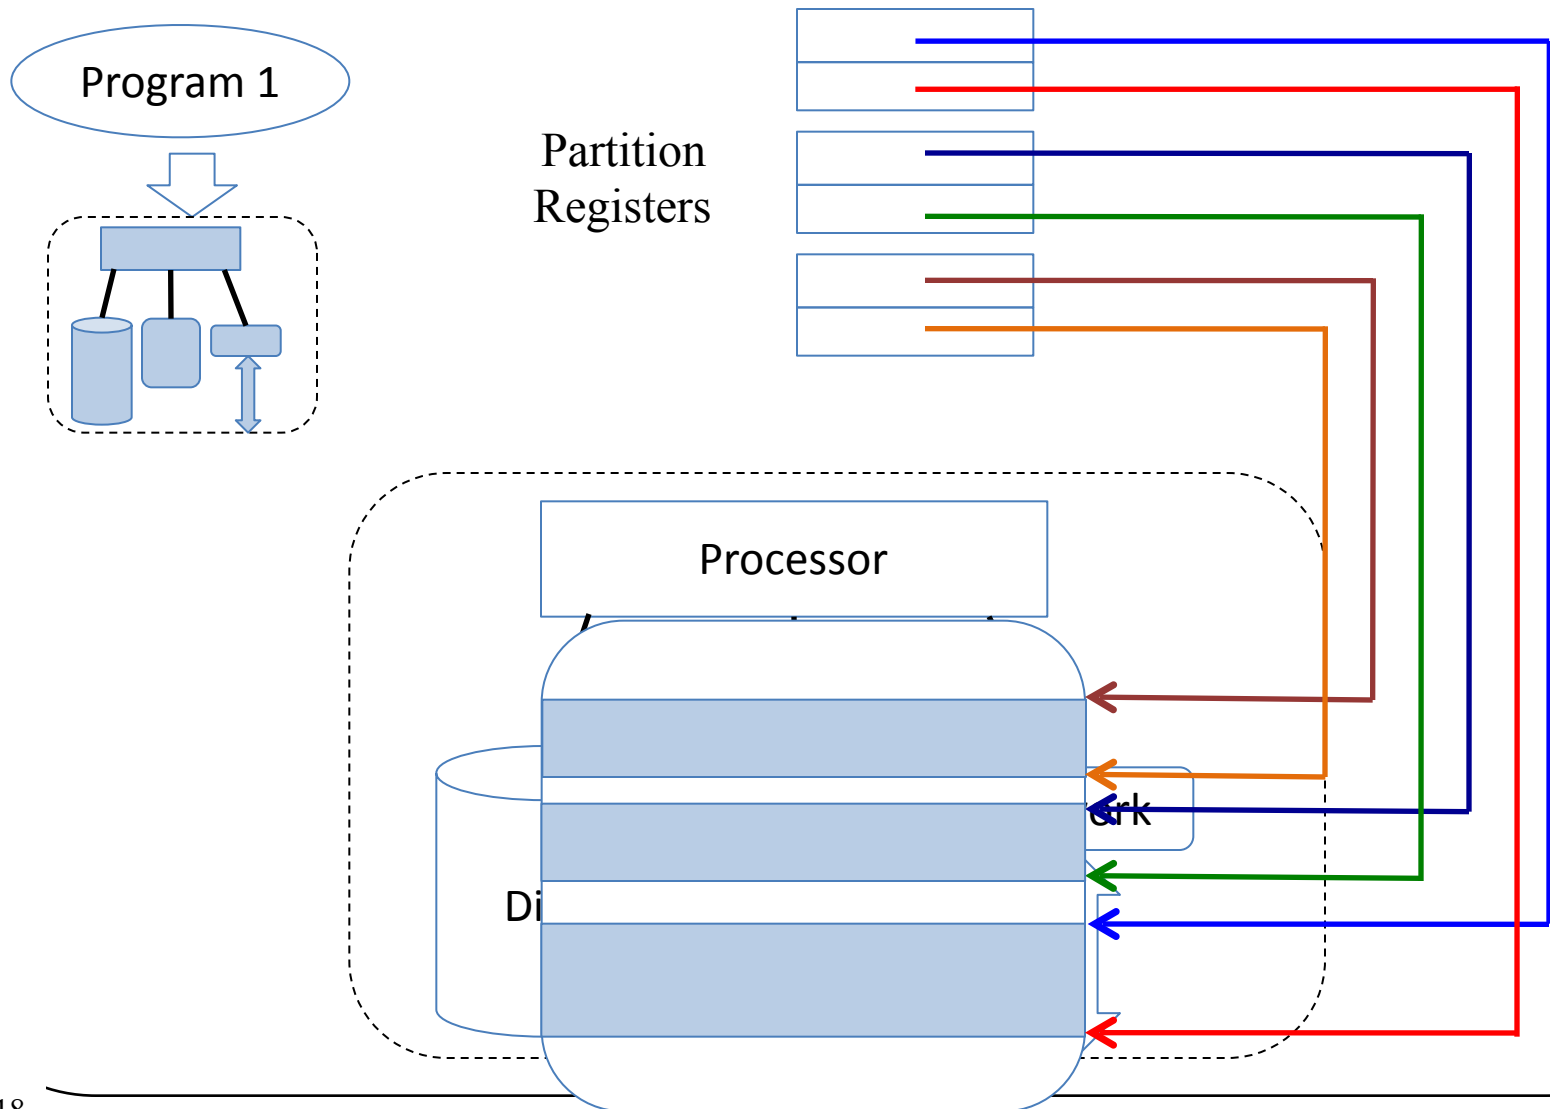
# Fixed Partition Allocation

- Pre-allocate partitions for  $n$  processes
  - One or more per process
  - Reserving space for largest possible process
- Partitions come in one or a few set sizes
- Very easy to implement
  - Common in old batch processing systems
  - Allocation/deallocation very cheap and easy
- Well suited to well-known job mix

# Memory Protection and Fixed Partitions

- Need to enforce partition boundaries
  - To prevent one process from accessing another's memory
- Could use hardware for this purpose
  - Special registers that contain the partition boundaries
  - Only accept addresses within the register values
- Basic scheme doesn't use virtual addresses

# The Partition Concept



# Problems With Fixed Partition Allocation

- Presumes you know how much memory will be used ahead of time
- Limits the number of processes supported to the total of their memory requirements
- Not great for sharing memory
- *Fragmentation* causes inefficient memory use

# Fragmentation

- A problem for all memory management systems
  - Fixed partitions suffer it especially badly
- Based on processes not using all the memory they requested
- As a result, you can't provide memory for as many processes as you theoretically could

# Fragmentation Example

Let's say there are three processes, A, B, and C

Their memory requirements:

A: 6 MBytes

B: 3 MBytes


C: 2 MBytes

Available partition sizes:


8 Mbytes

4 Mbytes


4 Mbytes



Partition 1  
8MB



Partition 2  
4MB



Partition 3  
4MB

$$\text{Total waste} = 2\text{MB} + 1\text{MB} + 2\text{MB} = 5/16\text{MB} = 31\%$$

# Internal Fragmentation

- Fragmentation comes in two kinds:
  - Internal and external
- This is an example of *internal fragmentation*
  - We'll see external fragmentation later
- Wasted space *inside* fixed sized blocks
  - The requestor was given more than he needed
  - The unused part is wasted and can't be used for others
- Internal fragmentation can occur whenever you force allocation in fixed-sized chunks

# More on Internal Fragmentation

- Internal fragmentation is caused by a mismatch between
  - The chosen size of a fixed-sized block
  - The actual sizes that programs use
- Average waste: 50% of each block



# Summary of Fixed Partition Allocation

- Very simple
- Inflexible
- Subject to a lot of internal fragmentation
- Not used in many modern systems
  - But a possible option for special purpose systems, like embedded systems
  - Where we know exactly what our memory needs will be

# Dynamic Partition Allocation

- Like fixed partitions, except
  - Variable sized, usually any size requested
  - Each partition has contiguous memory addresses
  - Processes have access permissions for the partitions
  - Potentially shared between processes
- Each process could have multiple partitions
  - With different sizes and characteristics
- In basic scheme, still only physical addresses

# Problems With Dynamic Partitions

- Not relocatable
  - Once a process has a partition, you can't easily move its contents elsewhere
- Not easily expandable
- Impossible to support applications with larger address spaces than physical memory
  - Also can't support several applications whose total needs are greater than physical memory
- Also subject to fragmentation


# Relocation and Expansion

- Partitions are tied to particular address ranges
  - At least during an execution
- Can't just move the contents of a partition to another set of addresses
  - All the pointers in the contents will be wrong
  - And generally you don't know which memory locations contain pointers
- Hard to expand because there may not be space “nearby”

# The Expansion Problem

- Partitions are allocated on request
- Processes may ask for new ones later
- But partitions that have been given can't be moved somewhere else in memory
- Memory management system might have allocated all the space after a given partition
  - In which case, it can't be expanded

# Illustrating the Problem



Now Process B wants to expand its partition size

But if we do that, Process B steps on Process C's memory

We can't move C's partition out of the way

And we can't move B's partition to a free area

We're stuck, and must deny an expansion request that we have enough memory to handle

# How To Keep Track of Variable Sized Partitions?

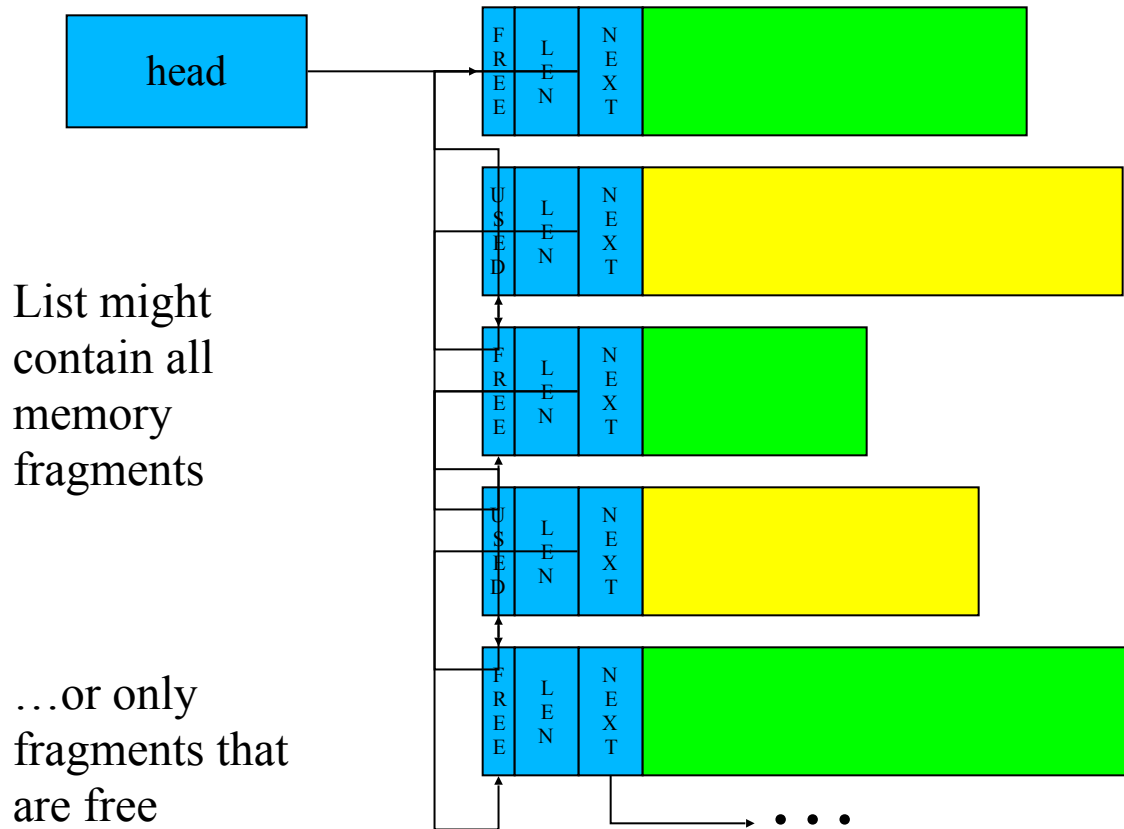
- Start with one large “heap” of memory
- Maintain a *free list*
  - Systems data structure to keep track of pieces of unallocated memory
- When a process requests more memory:
  - Find a large enough chunk of memory
  - Carve off a piece of the requested size
  - Put the remainder back on the free list
- When a process frees memory
  - Put it back on the free list

# Managing the Free List

- Fixed sized blocks are easy to track
  - A bit map indicating which blocks are free
- Variable chunks require more information
  - A linked list of descriptors, one per chunk
  - Each descriptor lists the size of the chunk and whether it is free
  - Each has a pointer to the next chunk on list
  - Descriptors often kept at front of each chunk
- Allocated memory may have descriptors too



# The Free List



# Free Chunk Carving


1. Find a large enough free chunk

2. Reduce its len to requested size

3. Create a new header for residual chunk

4. Insert the new chunk into the list


5. Mark the carved piece as in use



# Variable Partitions and Fragmentation

- Variable sized partitions not as subject to internal fragmentation
  - Unless requestor asked for more than he will use
  - Which is actually pretty common
  - But at least memory manager gave him no more than he requested
- Unlike fixed sized partitions, though, subject to another kind of fragmentation
  - *External fragmentation*

# External Fragmentation



We gradually build up small, unusable memory chunks scattered through memory


# External Fragmentation: Causes and Effects

- Each allocation creates left-over chunks
  - Over time they become smaller and smaller
- The small left-over fragments are useless
  - They are too small to satisfy any request
  - A second form of fragmentation waste
- Solutions:
  - Try not to create tiny fragments
  - Try to recombine fragments into big chunks

# How To Avoid Creating Small Fragments?

- Be smart about which free chunk of memory you use to satisfy a request
- But being smart costs time
- Some choices:
  - Best fit
  - Worst fit
  - First fit
  - Next fit

# Allocating Partitions in Memory




# Best Fit

- Search for the “best fit” chunk
  - Smallest size greater than or equal to requested size
- Advantages:
  - Might find a perfect fit
- Disadvantages:
  - Have to search entire list every time
  - Quickly creates very small fragments



# Best Fit in Action



# Worst Fit

- Search for the “worst fit” chunk
  - Largest size greater than or equal to requested size
- Advantages:
  - Tends to create very large fragments  
... for a while at least
- Disadvantages:
  - Still have to search entire list every time

# Worst Fit in Action



*Won't fit!*

*Won't fit!*


*Won't fit!*

# Comparing Best and Worst Fit

Best  
fit



Worst  
fit



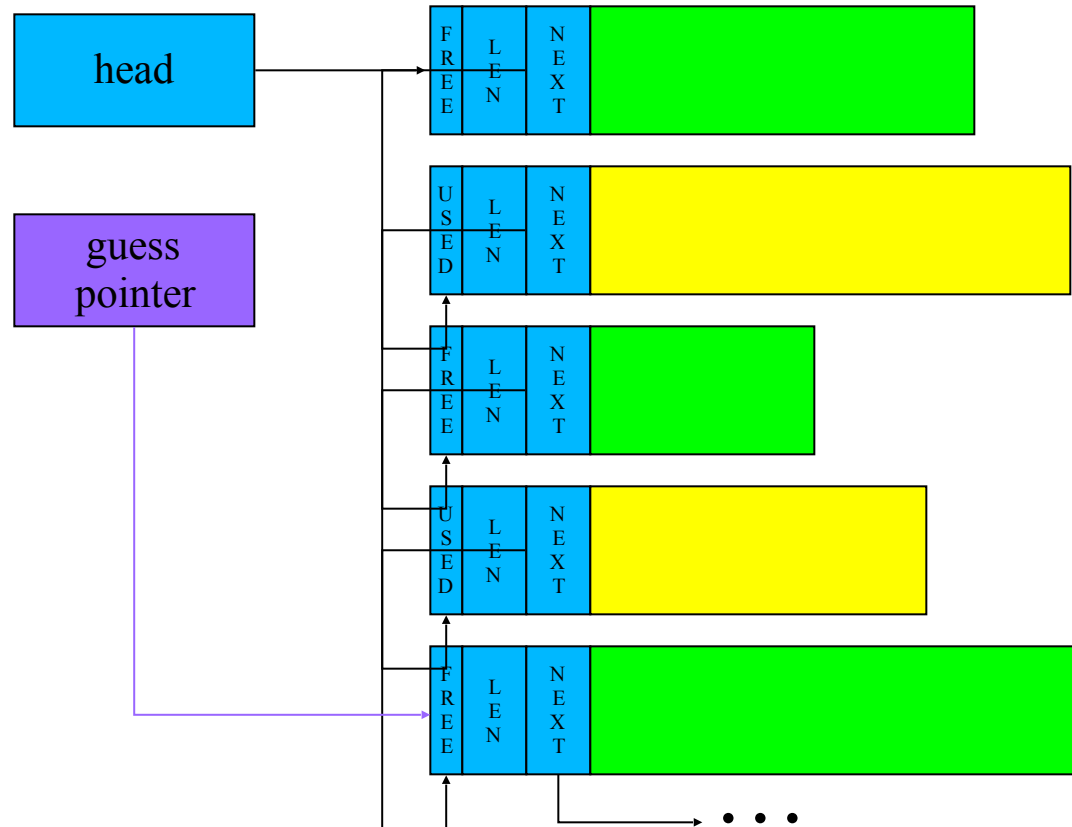
# First Fit

- Take first chunk you find that is big enough
- Advantages:
  - Very short searches
  - Creates random sized fragments
- Disadvantages:
  - The first chunks quickly fragment
  - Searches become longer
  - Ultimately it fragments as badly as best fit

# Next Fit

After each search, set guess pointer to chunk after the one we chose.

That is the point at which we will begin our next search.



# Next Fit Properties


- Tries to get advantages of both first and worst fit
  - Short searches (maybe shorter than first fit)
  - Spreads out fragmentation (like worst fit)
- Guess pointers are a general technique
  - Think of them as a lazy (non-coherent) cache
  - If they are right, they save a lot of time
  - If they are wrong, the algorithm still works
  - They can be used in a wide range of problems

# Coalescing Partitions

- All variable sized partition allocation algorithms have external fragmentation
  - Some get it faster, some spread it out
- We need a way to reassemble fragments
  - Check neighbors whenever a chunk is freed
  - Recombine free neighbors whenever possible
  - Free list can be designed to make this easier
    - E.g., where are the neighbors of this chunk?
- Counters forces of external fragmentation



# Free Chunk Coalescing




# Fragmentation and Coalescing

- Opposing processes that operate in parallel
  - Which of the two processes will dominate?
- What fraction of space is typically allocated?
  - Coalescing works better with more free space
- How fast is allocated memory turned over?
  - Chunks held for long time cannot be coalesced
- How variable are requested chunk sizes?
  - High variability increases fragmentation rate
- How long will the program execute?
  - Fragmentation, like rust, gets worse with time

# Variable Sized Partition Summary

- Eliminates internal fragmentation
  - Each chunk is custom-made for requestor
- Implementation is more expensive
  - Long searches of complex free lists
  - Carving and coalescing
- External fragmentation is inevitable
  - Coalescing can counteract the fragmentation
- Must we choose the lesser of two evils?

# A Special Case for Fixed Allocations



# Why Aren't Memory Request Sizes Randomly Distributed?

- In real systems, some sizes are requested much more often than others
- Many key services use fixed-size buffers
  - File systems (for disk I/O)
  - Network protocols (for packet assembly)
  - Standard request descriptors
- These account for much transient use
  - They are continuously allocated and freed
- OS might want to handle them specially

# Buffer Pools

- If there are popular sizes,
  - Reserve special pools of fixed size buffers
  - Satisfy matching requests from those pools
- Benefit: improved efficiency
  - Much simpler than variable partition allocation
    - Eliminates searching, carving, coalescing
  - Reduces (or eliminates) external fragmentation
- But we must know how much to reserve
  - Too little, and the buffer pool will become a bottleneck
  - Too much, and we will have a lot of unused buffer space
- Only satisfy perfectly matching requests
  - Otherwise, back to internal fragmentation

# How Are Buffer Pools Used?

- Process requests a piece of memory for a special purpose
  - E.g., to send a message
- System supplies one element from buffer pool
- Process uses it, completes, frees memory
  - Maybe explicitly
  - Maybe implicitly, based on how such buffers are used
    - E.g., sending the message will free the buffer “behind the process’ back” once the message is gone

# Dynamically Sizing Buffer Pools

- If we run low on fixed sized buffers
  - Get more memory from the free list
  - Carve it up into more fixed sized buffers
- If our free buffer list gets too large
  - Return some buffers to the free list
- If the free list gets dangerously low
  - Ask each major service with a buffer pool to return space
- This can be tuned by a few parameters:
  - Low space (need more) threshold
  - High space (have too much) threshold
  - Nominal allocation (what we free down to)
- Resulting system is highly adaptive to changing loads



# Lost Memory

- One problem with buffer pools is memory leaks
  - The process is done with the memory
  - But doesn't free it
- Also a problem when a process manages its own memory space
  - E.g., it allocates a big area and maintains its own free list
- Long running processes with memory leaks can waste huge amounts of memory

# Garbage Collection

- One solution to memory leaks
- Don't count on processes to release memory
- Monitor how much free memory we've got
- When we run low, start garbage collection
  - Search data space finding every object pointer
  - Note address/size of all accessible objects
  - Compute the compliment (what is inaccessible)
  - Add all inaccessible memory to the free list

# How Do We Find All Accessible Memory?

- Object oriented languages often enable this
  - All object references are tagged
  - All object descriptors include size information
- It is often possible for system resources
  - Where all possible references are known
    - (E.g., we know who has which files open)
- How about for the general case?

# General Garbage Collection

- Well, what would you need to do?
- Find all the pointers in allocated memory
- Determine “how much” each points to
- Determine what is and is not still pointed to
- Free what isn’t pointed to
- Why might that be difficult?



# Problems With General Garbage Collection

- A location in the data or stack segments might seem to contain addresses, but ...
  - Are they truly pointers, or might they be other data types whose values happen to resemble addresses?
  - If pointers, are they themselves still accessible?
  - We might be able to infer this (recursively) for pointers in dynamically allocated structures ...
  - But what about pointers in statically allocated (potentially global) areas?
- And how much is “pointed to,” one word or a million?

# Compaction and Relocation

- Garbage collection is just another way to free memory
  - Doesn't greatly help or hurt fragmentation
- Ongoing activity can starve coalescing
  - Chunks reallocated before neighbors become free
- We could stop accepting new allocations
  - But resulting convoy on memory manager would trash throughput
- We need a way to rearrange active memory
  - Re-pack all processes in one end of memory
  - Create one big chunk of free space at other end

# Memory Compaction



*Now let's compact!*

*An obvious improvement!*

# All This Requires Is Relocation . . .

- The ability to move a process
  - From region where it was initially loaded
  - Into a new and different region of memory
- What's so hard about that?
- All addresses in the program will be wrong
  - References in the code segment
    - Calls and branches to other parts of the code
    - References to variables in the data segment
  - Plus new pointers created during execution
    - That point into data and stack segments




# Why Is Relocation Hard?

```
int *foo;
```

**Before**

Let's move  
the  
partition!


**After**



What's going  
to happen the  
next time we  
access foo?

23000

33000




Of course, we copy the partition's contents when we move it


# The Relocation Problem

- It is not generally feasible to relocate a process
    - Maybe we could relocate references to code
      - If we kept the relocation information around
    - But how can we relocate references to data?
      - Pointer values may have been changed
      - New pointers may have been created
  - We could never find/fix all address references
    - Like the general case of garbage collection
  - Can we make processes location independent?
-


# Virtual Address Spaces



*Virtual* address space  
(as seen by process)



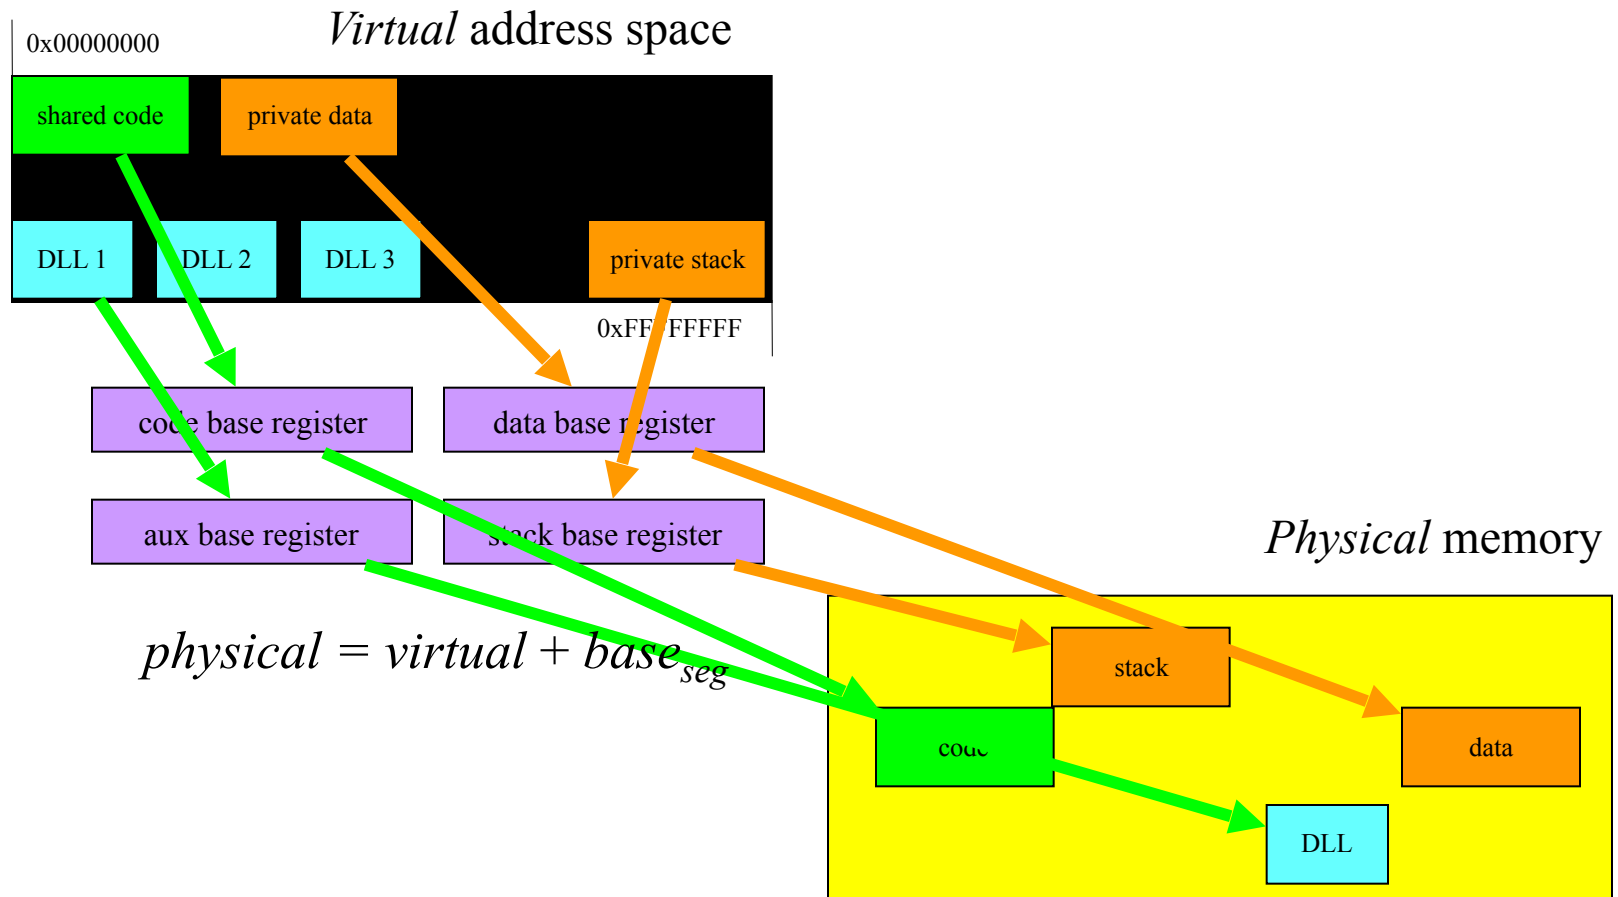
*Physical* address space  
(as on CPU/memory bus)



# Memory Segment Relocation

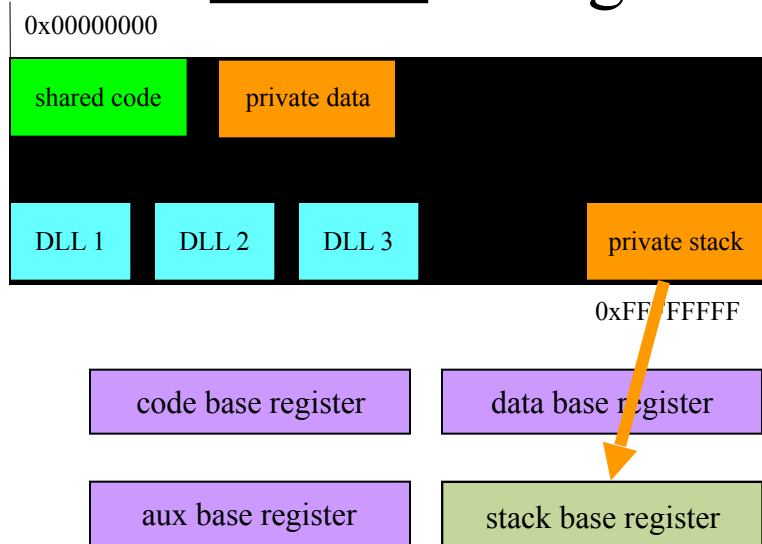
- A natural model
  - Process address space is made up of multiple segments
  - Use the segment as the unit of relocation
  - Long tradition, from the IBM system 360 to Intel x86 architecture
- Computer has special relocation registers
  - They are called segment base registers
  - They point to the start (in physical memory) of each segment
  - CPU automatically adds base register to every address
- OS uses these to perform virtual address translation
  - Set base register to start of region where program is loaded
  - If program is moved, reset base registers to new location
  - Program works no matter where its segments are loaded

# How Does Segment Relocation Work?



# Relocating a Segment

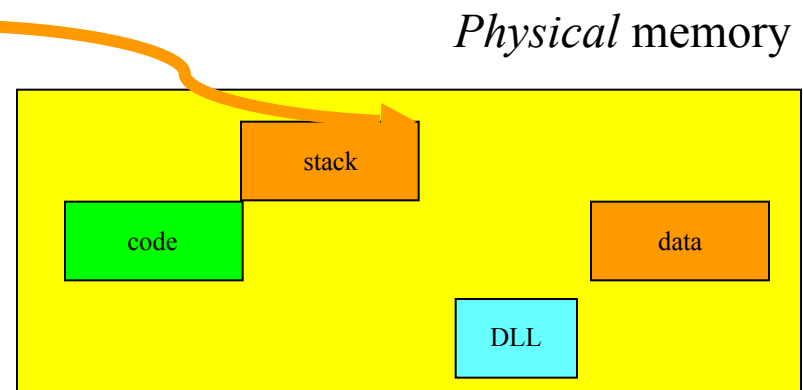
The virtual address of the stack doesn't change



Let's say we need to move the stack in physical memory

$$physical = virtual + base_{seg}$$

We just change the value in the stack base register



# Relocation and Safety

- A relocation mechanism (like base registers) is good
  - It solves the relocation problem
  - Enables us to move process segments in physical memory
  - Such relocation turns out to be insufficient
- We also need protection
  - Prevent process from reaching outside its allocated memory
    - E.g., by overrunning the end of a mapped segment
- Segments also need a length (or limit) register
  - Specifies maximum legal offset (from start of segment)
  - Any address greater than this is illegal (in the hole)
  - CPU should report it via a segmentation exception (trap)

# How Much of Our Problem Does Relocation Solve?

- We can use variable sized partitions
  - Cutting down on internal fragmentation
- We can move partitions around
  - Which helps coalescing be more effective
  - But still requires contiguous chunks of data for segments
  - So external fragmentation is still a problem
- We need to get rid of the requirement of contiguous segments